

Planar_data_classification_with_one_hidden_layer

May 18, 2025

1 Planar data classification with one hidden layer

Welcome to your week 3 programming assignment! It's time to build your first neural network, which will have one hidden layer. Now, you'll notice a big difference between this model and the one you implemented previously using logistic regression.

By the end of this assignment, you'll be able to:

- Implement a 2-class classification neural network with a single hidden layer
- Use units with a non-linear activation function, such as \tanh
- Compute the cross entropy loss
- Implement forward and backward propagation

1.1 Important Note on Submission to the AutoGrader

Before submitting your assignment to the AutoGrader, please make sure you are not doing the following:

1. You have not added any *extra* `print` statement(s) in the assignment.
2. You have not added any *extra* code cell(s) in the assignment.
3. You have not changed any of the function parameters.
4. You are not using any global variables inside your graded exercises. Unless specifically instructed to do so, please refrain from it and use the local variables instead.
5. You are not changing the assignment code where it is not required, like creating *extra* variables.

If you do any of the following, you will get something like, **Grader Error: Grader feedback not found** (or similarly unexpected) error upon submitting your assignment. Before asking for help/debugging the errors in your assignment, check for these first. If this is the case, and you don't remember the changes you have made, you can get a fresh copy of the assignment by following these [instructions](#).

1.2 Table of Contents

- Section ??
- Section ??
 - Section ??
- Section ??
- Section ??

- Section ??
 - * Section ??
- Section ??
 - * Section ??
- Section ??
 - * Section ??
- Section ??
 - * Section ??
- Section ??
 - * Section ??
- Section ??
 - * Section ??
- Section ??
 - Section ??
 - * Section ??
 - Section ??
- Section ??
- Section ??

1 - Packages

First import all the packages that you will need during this assignment.

- [numpy](#) is the fundamental package for scientific computing with Python.
- [sklearn](#) provides simple and efficient tools for data mining and data analysis.
- [matplotlib](#) is a library for plotting graphs in Python.
- `testCases` provides some test examples to assess the correctness of your functions
- `planar_utils` provide various useful functions used in this assignment

```
[ ]: ### v3.0
```

```
[ ]: # Package imports
import numpy as np
import copy
import matplotlib.pyplot as plt
from testCases_v2 import *
from public_tests import *
import sklearn
import sklearn.datasets
import sklearn.linear_model
from planar_utils import plot_decision_boundary, sigmoid, load_planar_dataset, ↵
↵load_extra_datasets

%matplotlib inline

%load_ext autoreload
%autoreload 2
```

2 - Load the Dataset

```
[ ]: X, Y = load_planar_dataset()
```

Visualize the dataset using matplotlib. The data looks like a “flower” with some red (label $y=0$) and some blue ($y=1$) points. Your goal is to build a model to fit this data. In other words, we want the classifier to define regions as either red or blue.

```
[ ]: # Visualize the data:
plt.scatter(X[0, :], X[1, :], c=Y, s=40, cmap=plt.cm.Spectral);
```

You have: - a numpy-array (matrix) X that contains your features (x_1, x_2) - a numpy-array (vector) Y that contains your labels (red:0, blue:1).

First, get a better sense of what your data is like.

Exercise 1

How many training examples do you have? In addition, what is the **shape** of the variables X and Y ?

Hint: How do you get the shape of a numpy array? ([help](#))

```
[3]: import numpy as np

# Example: Define X and Y
X = np.random.randn(100, 5) # Example data (100 samples, 5 features)
Y = np.random.randn(100, 1) # Example labels (100 samples, 1 label)

# ( 3 lines of code)
shape_X = X.shape
shape_Y = Y.shape
m = X.shape[1]

# Print the shapes and m
print('The shape of X is: ' + str(shape_X))
print('The shape of Y is: ' + str(shape_Y))
print('I have m = %d training examples!' % (m))
```

The shape of X is: (100, 5)
The shape of Y is: (100, 1)
I have $m = 5$ training examples!

Expected Output:

shape of X
(2, 400)
shape of Y
(1, 400)

m

400

3 - Simple Logistic Regression

Before building a full neural network, let's check how logistic regression performs on this problem. You can use sklearn's built-in functions for this. Run the code below to train a logistic regression classifier on the dataset.

```
[4]: import numpy as np
import sklearn
from sklearn.linear_model import LogisticRegressionCV

# Example dataset (replace with your actual data)
X = np.random.randn(2, 100) # 2 features, 100 samples
Y = np.random.randint(0, 2, (1, 100)) # Binary labels (0 or 1)

# Train the logistic regression classifier
clf = sklearn.linear_model.LogisticRegressionCV()
clf.fit(X.T, Y.T)

# Now the classifier 'clf' is trained on the data X.T and Y.T
```

```
[4]: LogisticRegressionCV(Cs=10, class_weight=None, cv=None, dual=False,
    fit_intercept=True, intercept_scaling=1.0, l1_ratios=None,
    max_iter=100, multi_class='auto', n_jobs=None,
    penalty='l2', random_state=None, refit=True, scoring=None,
    solver='lbfgs', tol=0.0001, verbose=0)
```

You can now plot the decision boundary of these models! Run the code below.

```
[6]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegressionCV

# Example dataset (replace with your actual data)
X = np.random.randn(2, 100) # 2 features, 100 samples
Y = np.random.randint(0, 2, (1, 100)) # Binary labels (0 or 1)

# Train the logistic regression classifier
clf = LogisticRegressionCV()
clf.fit(X.T, Y.T)

# Function to plot the decision boundary
def plot_decision_boundary(pred_func, X, Y):
    # Set up a grid for the plot
    x_min, x_max = X[0, :].min() - 1, X[0, :].max() + 1
    y_min, y_max = X[1, :].min() - 1, X[1, :].max() + 1
```

```

h = 0.02 # Step size for the grid


xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
Z = pred_func(np.c_[xx.ravel(), yy.ravel()]) # Predict labels for the grid
Z = Z.reshape(xx.shape) # Reshape to match the grid

# Plot the decision boundary and data points
plt.contourf(xx, yy, Z, alpha=0.8)
plt.scatter(X[0, :], X[1, :], c=Y.flatten(), edgecolors='k', marker='o')
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')

# Plot the decision boundary for logistic regression
plot_decision_boundary(lambda x: clf.predict(x), X, Y)
plt.title("Logistic Regression")
plt.show()

# Print accuracy
LR_predictions = clf.predict(X.T)
accuracy = float((np.dot(Y, LR_predictions) + np.dot(1 - Y, 1 - LR_predictions)) / float(Y.size)) * 100
print('Accuracy of logistic regression: %d%% ' % accuracy)

```



Accuracy of logistic regression: 54%

Expected Output:

Accuracy

47%

Interpretation: The dataset is not linearly separable, so logistic regression doesn't perform well. Hopefully a neural network will do better. Let's try this now!

4 - Neural Network model

Logistic regression didn't work well on the flower dataset. Next, you're going to train a Neural Network with a single hidden layer and see how that handles the same problem.

The model:

Mathematically:

For one example $x^{(i)}$:

$$z^{[1](i)} = W^{[1]}x^{(i)} + b^{[1]} \quad (1)$$

$$a^{[1](i)} = \tanh(z^{[1](i)}) \quad (2)$$

$$z^{[2](i)} = W^{[2]}a^{[1](i)} + b^{[2]} \quad (3)$$

$$\hat{y}^{(i)} = a^{[2](i)} = \sigma(z^{[2](i)}) \quad (4)$$

$$y_{prediction}^{(i)} = \begin{cases} 1 & \text{if } a^{[2](i)} > 0.5 \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

Given the predictions on all the examples, you can also compute the cost J as follows:

$$J = -\frac{1}{m} \sum_{i=0}^m \left(y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)}) \right) \quad (6)$$

Reminder: The general methodology to build a Neural Network is to: 1. Define the neural network structure (# of input units, # of hidden units, etc). 2. Initialize the model's parameters 3. Loop: - Implement forward propagation - Compute loss - Implement backward propagation to get the gradients - Update parameters (gradient descent)

In practice, you'll often build helper functions to compute steps 1-3, then merge them into one function called `nn_model()`. Once you've built `nn_model()` and learned the right parameters, you can make predictions on new data.

4.1 - Defining the neural network structure

Exercise 2 - layer_sizes

Define three variables: - `n_x`: the size of the input layer - `n_h`: the size of the hidden layer (**set this to 4, as `n_h = 4`, but only for this Exercise 2**) - `n_y`: the size of the output layer

Hint: Use shapes of `X` and `Y` to find `n_x` and `n_y`. Also, hard code the hidden layer size to be 4.

```
[15]: def layer_sizes(X, Y):
      """
      Arguments:
      X -- input dataset of shape (input size, number of examples)
      Y -- labels of shape (output size, number of examples)

      Returns:
      n_x -- the size of the input layer
      n_h -- the size of the hidden layer
      n_y -- the size of the output layer
      """
      # The size of the input layer is the number of features in X (number of
      →rows in X)
      n_x = X.shape[0] # number of features (input size)

      # The size of the hidden layer is hardcoded to 4 as specified in the problem
      n_h = 4 # hidden layer size (hardcoded)

      # The size of the output layer is the number of classes in Y (number of
      →rows in Y)
      n_y = Y.shape[0] # number of classes (output size)

      return (n_x, n_h, n_y)
```

```
[16]: t_X, t_Y = layer_sizes_test_case()
      (n_x, n_h, n_y) = layer_sizes(t_X, t_Y)
      print("The size of the input layer is: n_x = " + str(n_x))
      print("The size of the hidden layer is: n_h = " + str(n_h))
      print("The size of the output layer is: n_y = " + str(n_y))

      layer_sizes_test(layer_sizes)
```

```

      □
      →-----

      NameError                                Traceback (most recent call
      →last)

      <ipython-input-16-a625d2d64083> in <module>
      ----> 1 t_X, t_Y = layer_sizes_test_case()
            2 (n_x, n_h, n_y) = layer_sizes(t_X, t_Y)
            3 print("The size of the input layer is: n_x = " + str(n_x))
            4 print("The size of the hidden layer is: n_h = " + str(n_h))
            5 print("The size of the output layer is: n_y = " + str(n_y))
```

NameError: name 'layer_sizes_test_case' is not defined

Expected output

The size of the input layer is: $n_x = 5$
The size of the hidden layer is: $n_h = 4$
The size of the output layer is: $n_y = 2$
All tests passed!

4.2 - Initialize the model's parameters

Exercise 3 - initialize_parameters

Implement the function `initialize_parameters()`.

Instructions: - Make sure your parameters' sizes are right. Refer to the neural network figure above if needed. - You will initialize the weights matrices with random values. - Use: `np.random.randn(a,b) * 0.01` to randomly initialize a matrix of shape (a,b). - You will initialize the bias vectors as zeros. - Use: `np.zeros((a,b))` to initialize a matrix of shape (a,b) with zeros.

```
[28]: import numpy as np

def initialize_parameters(n_x, n_h, n_y):
    """
    Argument:
    n_x -- size of the input layer
    n_h -- size of the hidden layer
    n_y -- size of the output layer

    Returns:
    params -- python dictionary containing your parameters:
                W1 -- weight matrix of shape (n_h, n_x)
                b1 -- bias vector of shape (n_h, 1)
                W2 -- weight matrix of shape (n_y, n_h)
                b2 -- bias vector of shape (n_y, 1)
    """
    # Initialize W1, W2 with random small values, and b1, b2 as zeros
    W1 = np.random.randn(n_h, n_x) * 0.01 # shape: (n_h, n_x)
    b1 = np.zeros((n_h, 1)) # shape: (n_h, 1)
    W2 = np.random.randn(n_y, n_h) * 0.01 # shape: (n_y, n_h)
    b2 = np.zeros((n_y, 1)) # shape: (n_y, 1)

    # Store all parameters in a dictionary
    parameters = {"W1": W1,
                  "b1": b1,
                  "W2": W2,
                  "b2": b2}

    return parameters
```



```
[29]: np.random.seed(2)
n_x, n_h, n_y = initialize_parameters_test_case()
parameters = initialize_parameters(n_x, n_h, n_y)

print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
print("b2 = " + str(parameters["b2"]))

initialize_parameters_test(initialize_parameters)
```

```

↳ -----

NameError                                Traceback (most recent call↳
↳last)

<ipython-input-29-0eb4c3a6d62e> in <module>
      1 np.random.seed(2)
----> 2 n_x, n_h, n_y = initialize_parameters_test_case()
      3 parameters = initialize_parameters(n_x, n_h, n_y)
      4
      5 print("W1 = " + str(parameters["W1"]))

NameError: name 'initialize_parameters_test_case' is not defined
```

Expected output

```
W1 = [[-0.00416758 -0.00056267]
      [-0.02136196  0.01640271]
      [-0.01793436 -0.00841747]
      [ 0.00502881 -0.01245288]]
b1 = [[0.]
      [0.]
      [0.]
      [0.]]
W2 = [[-0.01057952 -0.00909008  0.00551454  0.02292208]]
b2 = [[0.]]
All tests passed!
```

4.3 - The Loop

Exercise 4 - forward_propagation

Implement `forward_propagation()` using the following equations:

$$Z^{[1]} = W^{[1]}X + b^{[1]} \quad (1)$$

$$A^{[1]} = \tanh(Z^{[1]}) \quad (2)$$

$$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]} \quad (3)$$

$$\hat{Y} = A^{[2]} = \sigma(Z^{[2]}) \quad (4)$$

Instructions:

- Check the mathematical representation of your classifier in the figure above.
- Use the function `sigmoid()`. It's built into (imported) this notebook.
- Use the function `np.tanh()`. It's part of the numpy library.
- Implement using these steps:
 1. Retrieve each parameter from the dictionary "parameters" (which is the output of `initialize_parameters()`) by using `parameters[".."]`.
 2. Implement Forward Propagation. Compute $Z^{[1]}$, $A^{[1]}$, $Z^{[2]}$ and $A^{[2]}$ (the vector of all your predictions on all the examples in the training set).
- Values needed in the backpropagation are stored in "cache". The cache will be given as an input to the backpropagation function.

```
[33]: import numpy as np

# Sigmoid function definition
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Forward propagation function
def forward_propagation(X, parameters):
    """
    Argument:
    X -- input data of size (n_x, m)
    parameters -- python dictionary containing your parameters (output of
    ↪ initialization function)

    Returns:
    A2 -- The sigmoid output of the second activation (probabilities)
    cache -- a dictionary containing "Z1", "A1", "Z2", and "A2"
    """

    # Retrieve each parameter from the dictionary "parameters"
    W1 = parameters["W1"]
    b1 = parameters["b1"]
    W2 = parameters["W2"]
    b2 = parameters["b2"]

    # Implement Forward Propagation to calculate A2 (probabilities)
    Z1 = np.dot(W1, X) + b1 # Z1 = W1 * X + b1
    A1 = np.tanh(Z1) # A1 = tanh(Z1)
    Z2 = np.dot(W2, A1) + b2 # Z2 = W2 * A1 + b2
```

```

A2 = sigmoid(Z2)  # A2 = sigmoid(Z2)

# Ensure that A2 has the correct shape (1, m)
assert(A2.shape == (1, X.shape[1]))

# Store the intermediate values in the cache dictionary
cache = {"Z1": Z1,
         "A1": A1,
         "Z2": Z2,
         "A2": A2}

return A2, cache

# Example test case
np.random.seed(2)

# Assuming forward_propagation_test_case is defined
t_X = np.random.randn(5, 100)  # Input data with 5 features, 100 samples
parameters = {
    "W1": np.random.randn(4, 5) * 0.01,  # W1: (4, 5)
    "b1": np.zeros((4, 1)),  # b1: (4, 1)
    "W2": np.random.randn(1, 4) * 0.01,  # W2: (1, 4)
    "b2": np.zeros((1, 1))  # b2: (1, 1)
}

```

```

[34]: t_X, parameters = forward_propagation_test_case()
      A2, cache = forward_propagation(t_X, parameters)
      print("A2 = " + str(A2))

      forward_propagation_test(forward_propagation)

```

```

↳ -----
NameError                                Traceback (most recent call↳
↳ last)

```

```

<ipython-input-34-e41f1de7d764> in <module>
----> 1 t_X, parameters = forward_propagation_test_case()
      2 A2, cache = forward_propagation(t_X, parameters)
      3 print("A2 = " + str(A2))
      4
      5 forward_propagation_test(forward_propagation)

```

```

NameError: name 'forward_propagation_test_case' is not defined

```

Expected output

```
A2 = [[0.21292656 0.21274673 0.21295976]]
```

```
All tests passed!
```

```
### 4.4 - Compute the Cost
```

Now that you've computed $A^{[2]}$ (in the Python variable "A2"), which contains $a^{[2](i)}$ for all examples, you can compute the cost function as follows:

$$J = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{[2](i)}) + (1 - y^{(i)}) \log(1 - a^{[2](i)})) \quad (13)$$

```
### Exercise 5 - compute_cost
```

Implement `compute_cost()` to compute the value of the cost J .

Instructions: - There are many ways to implement the cross-entropy loss. This is one way to implement one part of the equation without for loops: $-\sum_{i=1}^m y^{(i)} \log(a^{[2](i)})$:

```
logprobs = np.multiply(np.log(A2), Y)
```

```
cost = - np.sum(logprobs)
```

- Use that to build the whole expression of the cost function.

Notes:

- You can use either `np.multiply()` and then `np.sum()` or directly `np.dot()`.
- If you use `np.multiply` followed by `np.sum` the end result will be a type `float`, whereas if you use `np.dot`, the result will be a 2D numpy array.
- You can use `np.squeeze()` to remove redundant dimensions (in the case of single float, this will be reduced to a zero-dimension array).
- You can also cast the array as a type `float` using `float()`.

```
[35]: import numpy as np

def compute_cost(A2, Y):
    """
    Computes the cross-entropy cost given in equation (13)

    Arguments:
    A2 -- The sigmoid output of the second activation, of shape (1, number of
    → examples)
    Y -- "true" labels vector of shape (1, number of examples)

    Returns:
    cost -- cross-entropy cost given equation (13)
```

```

"""

m = Y.shape[1] # number of examples

# Compute the cross-entropy cost
logprobs = np.multiply(np.log(A2), Y) + np.multiply(np.log(1 - A2), (1 - Y))
cost = - np.sum(logprobs) / m

# Ensure cost is a scalar
cost = float(np.squeeze(cost)) # reduces dimensionality to scalar if
↳necessary

return cost

# Test the function with a test case
A2 = np.array([[0.8, 0.9, 0.1]]) # example predictions (probabilities)
t_Y = np.array([[1, 0, 1]]) # example true labels

cost = compute_cost(A2, t_Y)
print("cost = " + str(cost))

```

cost = 1.6094379124341003

```

[36]: A2, t_Y = compute_cost_test_case()
cost = compute_cost(A2, t_Y)
print("cost = " + str(compute_cost(A2, t_Y)))

compute_cost_test(compute_cost)

```

```

↳
-----

NameError                                Traceback (most recent call
↳last)

<ipython-input-36-9523b300d7a7> in <module>
----> 1 A2, t_Y = compute_cost_test_case()
      2 cost = compute_cost(A2, t_Y)
      3 print("cost = " + str(compute_cost(A2, t_Y)))
      4
      5 compute_cost_test(compute_cost)

NameError: name 'compute_cost_test_case' is not defined

```

Expected output

```
cost = 0.6930587610394646
```

```
All tests passed!
```

4.5 - Implement Backpropagation

Using the cache computed during forward propagation, you can now implement backward propagation.

Exercise 6 - backward_propagation

Implement the function `backward_propagation()`.

Instructions: Backpropagation is usually the hardest (most mathematical) part in deep learning. To help you, here again is the slide from the lecture on backpropagation. You'll want to use the six equations on the right of this slide, since you are building a vectorized implementation.

Figure 1: Backpropagation. Use the six equations on the right.

- Tips:
 - To compute $dZ1$ you'll need to compute $g^{[1]'}(Z^{[1]})$. Since $g^{[1]}(.)$ is the tanh activation function, if $a = g^{[1]}(z)$ then $g^{[1]'}(z) = 1 - a^2$. So you can compute $g^{[1]'}(Z^{[1]})$ using `(1 - np.power(A1, 2))`.

```
[38]: def backward_propagation(parameters, cache, X, Y):
    """
    Implement the backward propagation using the instructions above.

    Arguments:
    parameters -- python dictionary containing our parameters
    cache -- a dictionary containing "Z1", "A1", "Z2" and "A2".
    X -- input data of shape (n_x, number of examples)
    Y -- "true" labels vector of shape (1, number of examples)

    Returns:
    grads -- python dictionary containing your gradients with respect to
    ↪ different parameters
    """
    m = X.shape[1]

    # Retrieve W1 and W2 from the dictionary "parameters".
    W1 = parameters["W1"]
    W2 = parameters["W2"]

    # Retrieve A1 and A2 from dictionary "cache".
    A1 = cache["A1"]
    A2 = cache["A2"]

    # Backward propagation: calculate dW1, db1, dW2, db2.
    dZ2 = A2 - Y # derivative of cost with respect to A2
    dW2 = (1/m) * np.dot(dZ2, A1.T) # gradient for W2
    db2 = (1/m) * np.sum(dZ2, axis=1, keepdims=True) # gradient for b2
```

```

    dZ1 = np.dot(W2.T, dZ2) * (1 - np.power(A1, 2)) # derivative of tanh
    ↪ activation
    dW1 = (1/m) * np.dot(dZ1, X.T) # gradient for W1
    db1 = (1/m) * np.sum(dZ1, axis=1, keepdims=True) # gradient for b1

    grads = {"dW1": dW1,
             "db1": db1,
             "dW2": dW2,
             "db2": db2}

    return grads

```

```

[ ]: parameters, cache, t_X, t_Y = backward_propagation_test_case()

grads = backward_propagation(parameters, cache, t_X, t_Y)
print ("dW1 = " + str(grads["dW1"]))
print ("db1 = " + str(grads["db1"]))
print ("dW2 = " + str(grads["dW2"]))
print ("db2 = " + str(grads["db2"]))

backward_propagation_test(backward_propagation)

```

Expected output

```

dW1 = [[ 0.00301023 -0.00747267]
       [ 0.00257968 -0.00641288]
       [-0.00156892  0.003893   ]
       [-0.00652037  0.01618243]]
db1 = [[ 0.00176201]
       [ 0.00150995]
       [-0.00091736]
       [-0.00381422]]
dW2 = [[ 0.00078841  0.01765429 -0.00084166 -0.01022527]]
db2 = [[-0.16655712]]
All tests passed!

```

4.6 - Update Parameters

Exercise 7 - update_parameters

Implement the update rule. Use gradient descent. You have to use (dW1, db1, dW2, db2) in order to update (W1, b1, W2, b2).

General gradient descent rule: $\theta = \theta - \alpha \frac{\partial J}{\partial \theta}$ where α is the learning rate and θ represents a parameter.

Figure 2: The gradient descent algorithm with a good learning rate (converging) and a bad learning rate (diverging). Images courtesy of Adam Harley.

Hint

- Use `copy.deepcopy(...)` when copying lists or dictionaries that are passed as parameters to functions. It avoids input parameters being modified within the function. In some scenarios, this could be inefficient, but it is required for grading purposes.

```
[41]: import copy

def update_parameters(parameters, grads, learning_rate = 1.2):
    """
    Updates parameters using the gradient descent update rule given above

    Arguments:
    parameters -- python dictionary containing your parameters
    grads -- python dictionary containing your gradients

    Returns:
    parameters -- python dictionary containing your updated parameters
    """
    # Retrieve a copy of each parameter from the dictionary "parameters".
    # Use copy.deepcopy(...) to avoid modifying the original parameters.
    W1 = copy.deepcopy(parameters["W1"])
    b1 = copy.deepcopy(parameters["b1"])
    W2 = copy.deepcopy(parameters["W2"])
    b2 = copy.deepcopy(parameters["b2"])

    # Retrieve each gradient from the dictionary "grads"
    dW1 = grads["dW1"]
    db1 = grads["db1"]
    dW2 = grads["dW2"]
    db2 = grads["db2"]

    # Update rule for each parameter
    W1 = W1 - learning_rate * dW1
    b1 = b1 - learning_rate * db1
    W2 = W2 - learning_rate * dW2
    b2 = b2 - learning_rate * db2

    # Store the updated parameters in the dictionary
    parameters = {"W1": W1,
                  "b1": b1,
                  "W2": W2,
                  "b2": b2}

    return parameters
```

```
[42]: parameters, grads = update_parameters_test_case()
parameters = update_parameters(parameters, grads)
```



```

print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
print("b2 = " + str(parameters["b2"]))

update_parameters_test(update_parameters)

```

```

↳ -----

NameError                                Traceback (most recent call↳
↳ last)

<ipython-input-42-c37f65e5c916> in <module>
----> 1 parameters, grads = update_parameters_test_case()
      2 parameters = update_parameters(parameters, grads)
      3
      4 print("W1 = " + str(parameters["W1"]))
      5 print("b1 = " + str(parameters["b1"]))

```

NameError: name 'update_parameters_test_case' is not defined

Expected output

```

W1 = [[-0.00643025  0.01936718]
      [-0.02410458  0.03978052]
      [-0.01653973 -0.02096177]
      [ 0.01046864 -0.05990141]]
b1 = [[-1.02420756e-06]
      [ 1.27373948e-05]
      [ 8.32996807e-07]
      [-3.20136836e-06]]
W2 = [[-0.01041081 -0.04463285  0.01758031  0.04747113]]
b2 = [[0.00010457]]
All tests passed!

```

4.7 - Integration

Integrate your functions in `nn_model()`

Exercise 8 - `nn_model`

Build your neural network model in `nn_model()`.

Instructions: The neural network model has to use the previous functions in the right order.

```
[43]: def nn_model(X, Y, n_h, num_iterations = 10000, print_cost=False):
    """
    Arguments:
    X -- dataset of shape (2, number of examples)
    Y -- labels of shape (1, number of examples)
    n_h -- size of the hidden layer
    num_iterations -- Number of iterations in gradient descent loop
    print_cost -- if True, print the cost every 1000 iterations

    Returns:
    parameters -- parameters learned by the model. They can then be used to
    ↪predict.
    """

    np.random.seed(3)

    # Get the size of the input and output layers
    n_x = X.shape[0] # number of features in the input
    n_y = Y.shape[0] # number of classes (output layer size)

    # Initialize parameters
    parameters = initialize_parameters(n_x, n_h, n_y)

    # Loop (gradient descent)
    for i in range(0, num_iterations):

        # Forward propagation. Inputs: "X, parameters". Outputs: "A2, cache".
        A2, cache = forward_propagation(X, parameters)

        # Compute the cost. Inputs: "A2, Y". Outputs: "cost".
        cost = compute_cost(A2, Y)

        # Backpropagation. Inputs: "parameters, cache, X, Y". Outputs: "grads".
        grads = backward_propagation(parameters, cache, X, Y)

        # Gradient descent parameter update. Inputs: "parameters, grads".
        ↪Outputs: "parameters".
        parameters = update_parameters(parameters, grads)

        # Print the cost every 1000 iterations
        if print_cost and i % 1000 == 0:
            print(f"Cost after iteration {i}: {cost}")

    return parameters
```

```
[ ]: nn_model_test(nn_model)
```

Expected output

```
Cost after iteration 0: 0.693497
Cost after iteration 1000: 0.000180
Cost after iteration 2000: 0.000088
...
Cost after iteration 8000: 0.000022
Cost after iteration 9000: 0.000019
W1 = [[-0.48394893 -0.91443482]
      [-0.69227123 -1.28596651]
      [ 0.63806018  1.18124948]
      [ 0.73594679  1.35718308]
      [-0.62621054 -1.16022636]]
b1 = [[ 0.00285386]
      [ 0.01642488]
      [-0.01114986]
      [-0.01656405]
      [ 0.01042274]]
W2 = [[-1.45325879 -2.65768451  2.26286606  2.97274518 -2.18860534]]
b2 = [[0.00758617]]
All tests passed!
```

5 - Test the Model

5.1 - Predict

Exercise 9 - predict

Predict with your model by building `predict()`. Use forward propagation to predict results.

Reminder: $\text{predictions} = y_{\text{prediction}} = \mathbb{I}(\text{activation} > 0.5) = \begin{cases} 1 & \text{if } \text{activation} > 0.5 \\ 0 & \text{otherwise} \end{cases}$

As an example, if you would like to set the entries of a matrix `X` to 0 and 1 based on a threshold you would do: `X_new = (X > threshold)`

```
[44]: def predict(parameters, X):
      """
      Using the learned parameters, predicts a class for each example in X

      Arguments:
      parameters -- python dictionary containing your parameters
      X -- input data of size (n_x, m)

      Returns:
      predictions -- vector of predictions of our model (red: 0 / blue: 1)
      """

      # Compute probabilities using forward propagation, and classify to 0/1
      → using 0.5 as the threshold.
      A2, cache = forward_propagation(X, parameters) # Forward propagation
```

```

predictions = (A2 > 0.5) # Apply threshold to classify

return predictions

```

```

[46]: parameters, t_X = predict_test_case()

predictions = predict(parameters, t_X)
print("Predictions: " + str(predictions))

predict_test(predict)

```

```

↳ -----

NameError                                Traceback (most recent call↳
↳last)

<ipython-input-46-eb08347ef392> in <module>
----> 1 parameters, t_X = predict_test_case()
      2
      3 predictions = predict(parameters, t_X)
      4 print("Predictions: " + str(predictions))
      5

NameError: name 'predict_test_case' is not defined

```

Expected output

```

Predictions: [[ True False  True]]
All tests passed!

```

5.2 - Test the Model on the Planar Dataset

It's time to run the model and see how it performs on a planar dataset. Run the following code to test your model with a single hidden layer of n_h hidden units!

```

[47]: # Run the model
parameters = nn_model(X, Y, n_h = 4, num_iterations = 10000, print_cost=True)

# Plot the decision boundary
plot_decision_boundary(lambda x: predict(parameters, x.T), X, Y)
plt.title("Decision Boundary for hidden layer size " + str(4))

# Calculate accuracy
predictions = predict(parameters, X)


```

```

accuracy = float((np.dot(Y, predictions.T) + np.dot(1 - Y, 1 - predictions.T)) /
    ↪ float(Y.size) * 100)
print ('Accuracy: %d' % accuracy + '%')

```

Cost after iteration 0: 0.6931479153847042
 Cost after iteration 1000: 0.6897854476294983
 Cost after iteration 2000: 0.6813674235618759
 Cost after iteration 3000: 0.6205054699797357
 Cost after iteration 4000: 0.5765689633126523
 Cost after iteration 5000: 0.5597015798602061
 Cost after iteration 6000: 0.5510388983668821
 Cost after iteration 7000: 0.5455540562200053
 Cost after iteration 8000: 0.5417411326638144
 Cost after iteration 9000: 0.5388463078510523
 Accuracy: 63%



```

[48]: # Print accuracy
predictions = predict(parameters, X)
print ('Accuracy: %d' % float((np.dot(Y, predictions.T) + np.dot(1 - Y, 1 -
    ↪ predictions.T)) / float(Y.size) * 100) + '%')

```

Accuracy: 63%

Expected Output:

Accuracy

90%

Accuracy is really high compared to Logistic Regression. The model has learned the patterns of the flower's petals! Unlike logistic regression, neural networks are able to learn even highly non-linear decision boundaries.

1.2.1 Congrats on finishing this Programming Assignment!

Here's a quick recap of all you just accomplished:

- Built a complete 2-class classification neural network with a hidden layer
- Made good use of a non-linear unit
- Computed the cross entropy loss
- Implemented forward and backward propagation
- Seen the impact of varying the hidden layer size, including overfitting.

You've created a neural network that can learn patterns! Excellent work. Below, there are some optional exercises to try out some other hidden layer sizes, and other datasets.

6 - Tuning hidden layer size (optional/ungraded exercise)

Run the following code(it may take 1-2 minutes). Then, observe different behaviors of the model for various hidden layer sizes.

```
[49]: # This may take about 2 minutes to run

plt.figure(figsize=(16, 32))
hidden_layer_sizes = [1, 2, 3, 4, 5]

# you can try with different hidden layer sizes
# but make sure before you submit the assignment it is set as
↪ "hidden_layer_sizes = [1, 2, 3, 4, 5]"
# hidden_layer_sizes = [1, 2, 3, 4, 5, 20, 50]


for i, n_h in enumerate(hidden_layer_sizes):
    plt.subplot(5, 2, i+1)
    plt.title('Hidden Layer of size %d' % n_h)
    parameters = nn_model(X, Y, n_h, num_iterations = 5000)
    plot_decision_boundary(lambda x: predict(parameters, x.T), X, Y)
    predictions = predict(parameters, X)
    accuracy = float((np.dot(Y,predictions.T) + np.dot(1 - Y, 1 - predictions.
↪ T)) / float(Y.size)*100)
    print ("Accuracy for {} hidden units: {} %".format(n_h, accuracy))
```

Accuracy for 1 hidden units: 56.99999999999999 %

Accuracy for 2 hidden units: 63.0 %

Accuracy for 3 hidden units: 64.0 %

Accuracy for 4 hidden units: 61.0 %
Accuracy for 5 hidden units: 69.0 %



Interpretation: - The larger models (with more hidden units) are able to fit the training set better, until eventually the largest models overfit the data. - The best hidden layer size seems to be around $n_h = 5$. Indeed, a value around here seems to fit the data well without also incurring noticeable overfitting. - Later, you'll become familiar with regularization, which lets you use very large models (such as $n_h = 50$) without much overfitting.

Note: Remember to submit the assignment by clicking the blue “Submit Assignment” button at the upper-right.

Some optional/ungraded questions that you can explore if you wish: - What happens when you change the tanh activation for a sigmoid activation or a ReLU activation? - Play with the learning_rate. What happens? - What if we change the dataset? (See part 7 below!)

7- Performance on other datasets

If you want, you can rerun the whole notebook (minus the dataset part) for each of the following datasets.

```
[ ]: # Datasets
noisy_circles, noisy_moons, blobs, gaussian_quantiles, no_structure = load_extra_datasets()

datasets = {"noisy_circles": noisy_circles,
            "noisy_moons": noisy_moons,
            "blobs": blobs,
            "gaussian_quantiles": gaussian_quantiles}

### START CODE HERE ### (choose your dataset)
dataset = "noisy_moons"
### END CODE HERE ###

X, Y = datasets[dataset]
X, Y = X.T, Y.reshape(1, Y.shape[0])

# make blobs binary
if dataset == "blobs":
    Y = Y%2

# Visualize the data
plt.scatter(X[0, :], X[1, :], c=Y, s=40, cmap=plt.cm.Spectral);
```

References:

- <http://cs231n.github.io/neural-networks-case-study/>

```
[ ]:
```